

Tire Condition Detection with a CNN + Streamlit App

3-class image classification: flat • full • no-tire | TensorFlow/Keras deployment

[Dataset → Training → Web Demo](#)

Why detect tire conditions from images?

Problem & Motivation



Real-world checks are hard

Manual inspections are slow; sensor solutions add hardware + maintenance. Vision-based checks can scale from simple cameras.



Safety impact

Flat or missing tires reduce stability and braking performance → higher risk incidents.



Cost & coverage

Cameras are cheaper and easier to deploy than specialized sensors (especially at scale).



ML opportunity

CNNs learn robust visual features (edges → textures → parts) directly from pixels, outperforming brittle hand-crafted rules in varied lighting/backgrounds.

Goal, Scope, and Deliverables

Goal

- Classify a user-uploaded RGB image into one of 3 classes: flat, full, or no-tire.
- Return a predicted label and confidence score (softmax probability).
- Deploy as a lightweight web app for interactive demos.

Deliverables

- Trained CNN model (.h5)
- Training curves (loss/accuracy)
- Streamlit inference UI
- Reproducible config (YAML)

Output example:

Prediction: full

Confidence: 89.44%

Dataset

Classes and split

- 3 classes: flat, full, no-tire
- Total images: 900
- Split: 80% train (720), 20% val (180)

Directory structure

```
data/raw/  
flat/  
full/  
no-tire/
```



Why a hold-out validation set?

Validation estimates generalization to unseen images and guides model selection (e.g., best epoch).



Why 3 classes (flat/full/no-tire)?

Covers the most actionable states for a demo: under-inflation, normal, and missing tire (data sanity).



Example tire image (visual variability)

Preprocessing & Augmentation

Pipeline

- Resize to 240×240 to standardize input and control compute.
- Normalize pixels with rescale = 1/255 to stabilize training.
- Augment only the training split: rotation ($\pm 10^\circ$), zoom ($\pm 10\%$), horizontal flip.
- Keep validation un-augmented to measure true generalization.

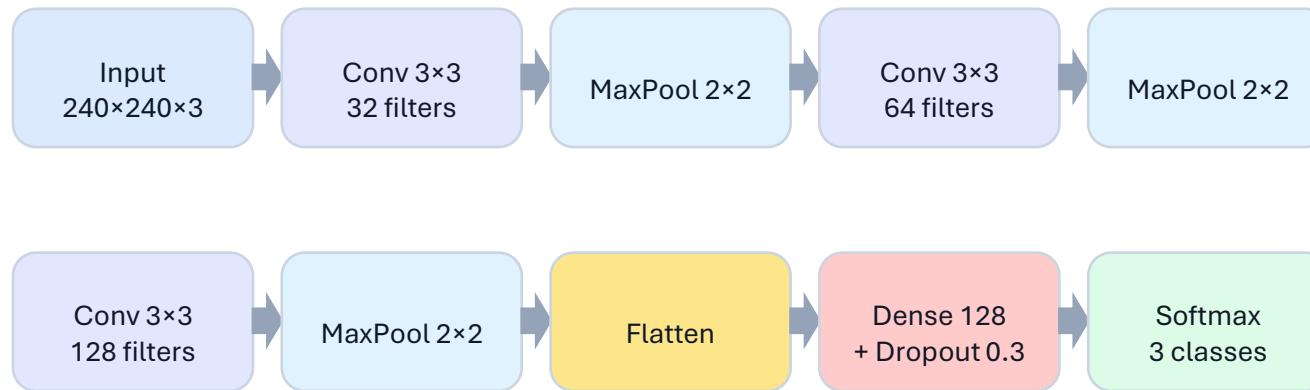
Reasoning for these choices

- Small geometric transforms simulate viewpoint variance
- Flip helps when left/right orientation is irrelevant
- Mild ranges reduce risk of generating unrealistic tires
- Data augmentation is a practical regularization tool

Implementation: Keras ImageDataGenerator

Model Architecture

Feature extractor → classifier head



Reasoning for choosing this CNN

- 1) Baseline-first: easy to train & debug
- 2) 3 conv blocks capture increasing complexity
- 3) Small 3x3 kernels are standard for locality
- 4) Dropout reduces overfitting with limited data
- 5) Softmax fits multi-class output

Forward Pass: Convolution

Convolution (with activation)

$$Y_{\{i,j,k\}} = \sigma(\sum_{u=1..r} \sum_{v=1..s} \sum_{c=1..C} W_{\{u,v,c,k\}} \cdot X_{\{i+u, j+v, c\}} + b_k)$$

Interpretation

- X is the input feature map (pixels or previous layer activations).
- W is the learnable kernel (filter) weights; b_k is the bias for filter k .
- The triple sum aggregates local neighborhoods across channels.
- σ is ReLU in this project: $\sigma(x)=\max(0,x)$.

Why it works for tires

- Early filters learn edges/curves (wheel boundary).
- Deeper filters learn textures (tread) and deformation cues (flat).
- Pooling adds translation tolerance (position shifts).

Decision Rule & Training Objective

Softmax → class probability

$$p_i = \exp(z_i) / \sum_{j=1..3} \exp(z_j)$$

$$\hat{y} = \operatorname{argmax}_i (p_i) ; \quad \text{confidence} = \max_i(p_i)$$

Categorical cross-entropy loss

$$L = - \sum_{i=1..3} y_i \cdot \log(p_i)$$

Minimizing L encourages high probability on the true class.



How this maps to your code

In Streamlit: `pred = model.predict(tensor)[0]` returns softmax probabilities.

`idx = argmax(pred)` selects the class; `confidence = pred[idx]`.

During training: `model.compile(loss="categorical_crossentropy", optimizer="adam")` minimizes cross-entropy over batches.

Training Setup



Core hyperparameters

Image size: 240×240

Batch size: 32

Epochs: 10

Classes: 3

Split: 80/20 via validation_split



Why Adam + Dropout?

Adam adapts learning rates per-parameter → strong default for CNN training.

Dropout (0.3) reduces co-adaptation and helps when Dense layer is large relative to dataset size.

Practical tip: select best epoch by validation accuracy (early stopping).

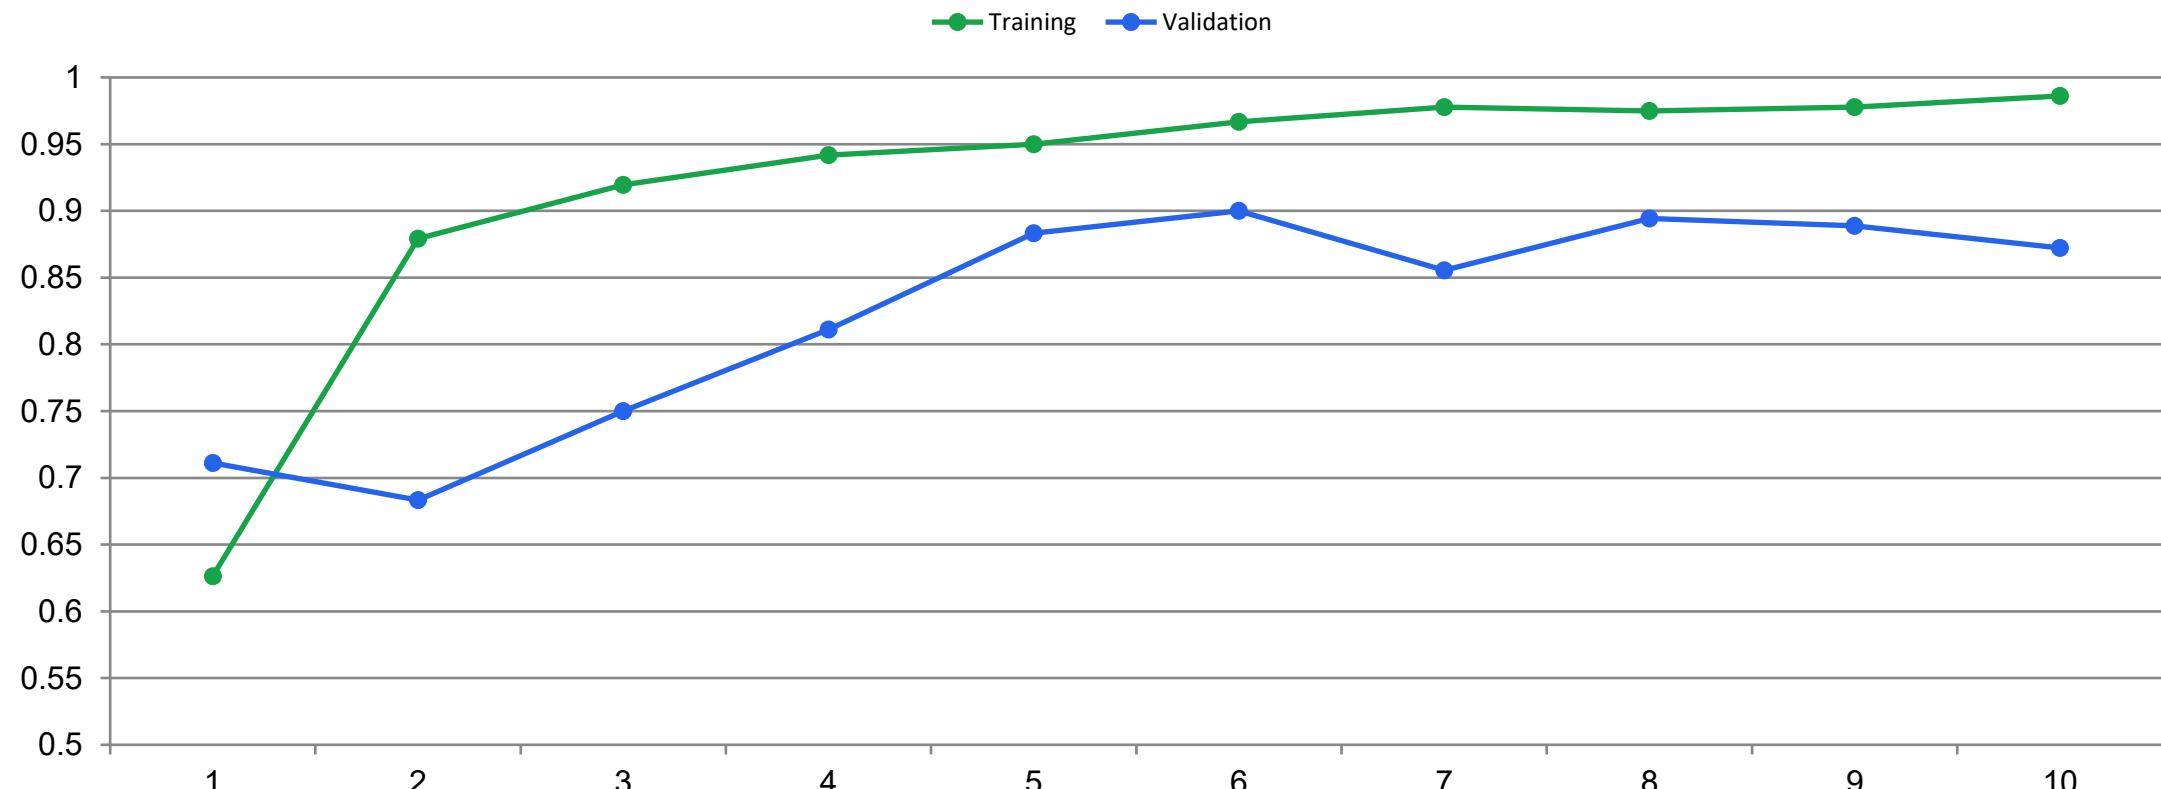


TensorFlow

Implementation stack

- TensorFlow / Keras (model, training)
- ImageDataGenerator (augmentation)
- Streamlit (web UI)
- YAML config (paths & class names)

Results



Best validation accuracy

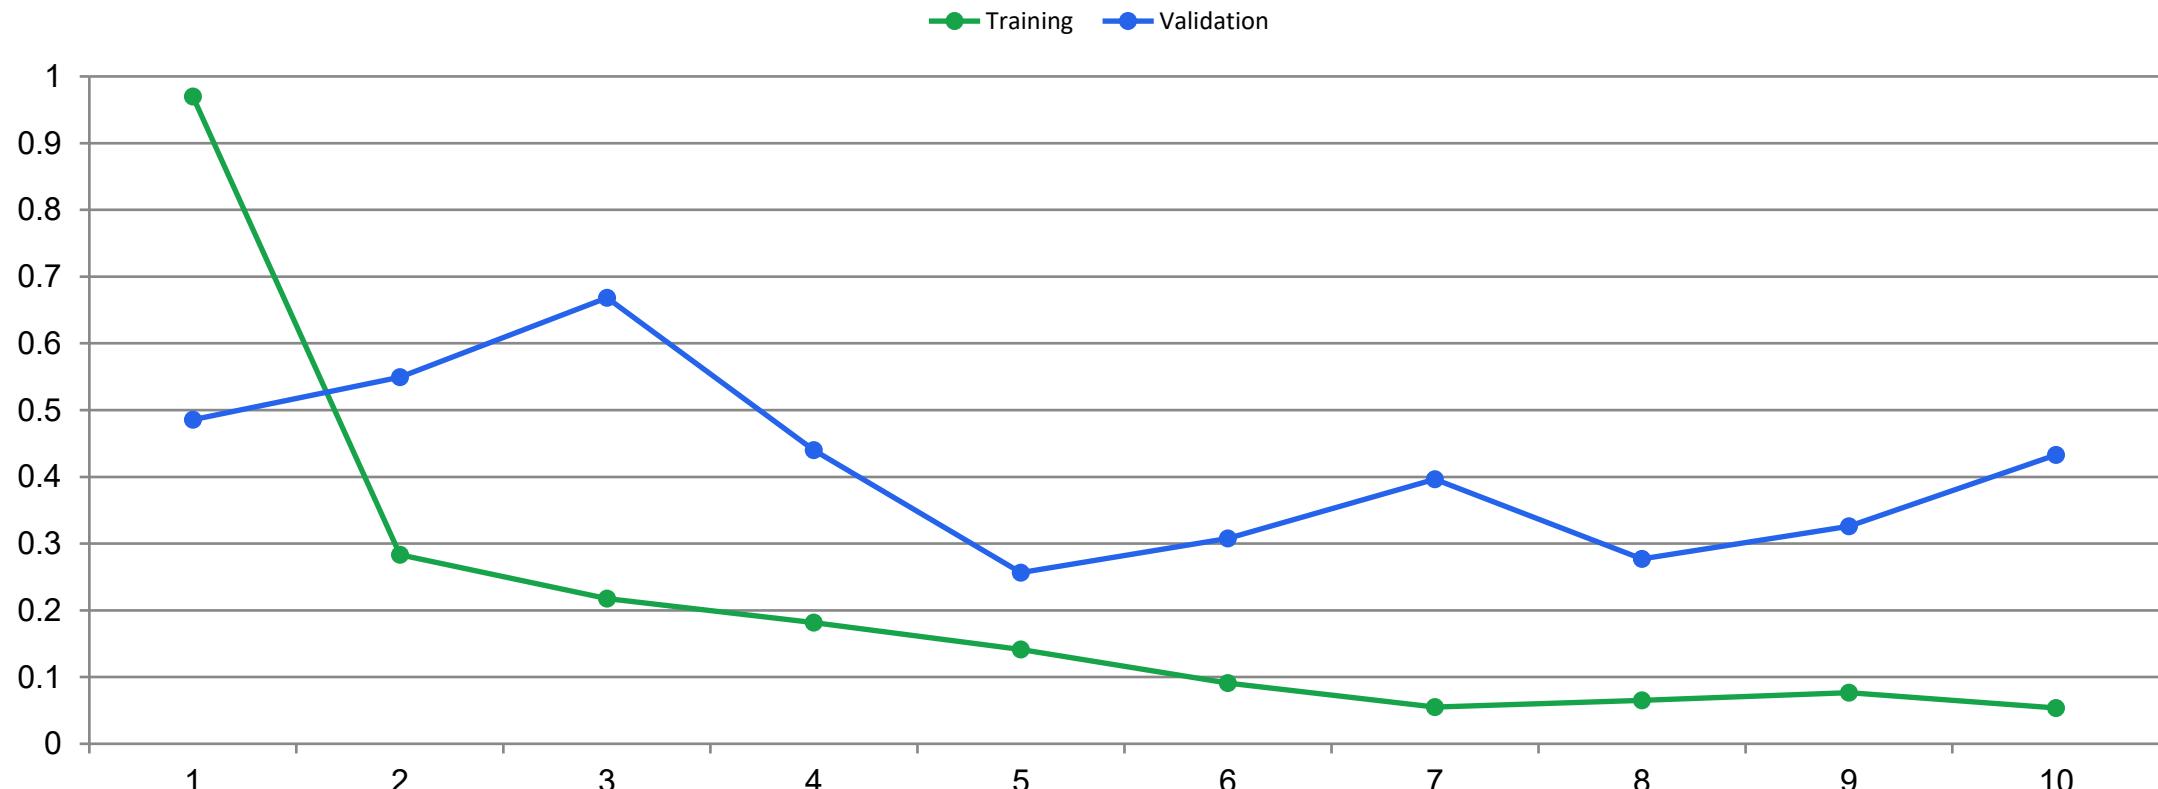
Peak val_acc ≈ 0.90 at Epoch 6 (candidate checkpoint).



Reading the curve

Training rises steadily; validation fluctuates → mild

Results

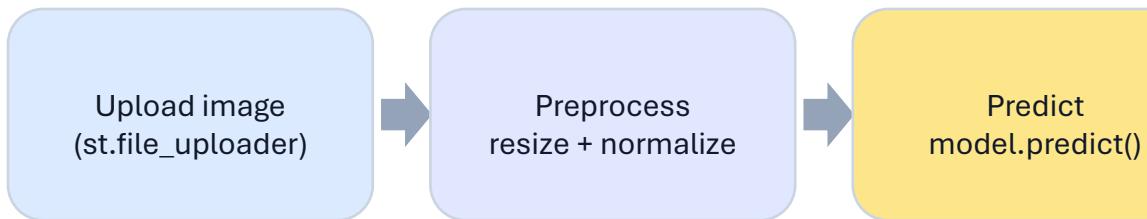


Interpretation

Training loss keeps dropping; validation loss bottoms out around Epoch 5–6 then rises → consider early stopping or stronger regularization.

Deployment

Inference flow



Key implementation choices

- YAML config keeps paths/class names centralized
- st.cache_resource loads the model once (fast reruns)
- preprocess_image mirrors training normalization
- Output: label + confidence (softmax)

```
@st.cache_resource
def load_model():
    return tf.keras.models.load_model(path)

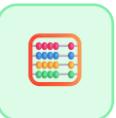
img = Image.open(uploaded).convert("RGB")
tensor = preprocess_image(img, IMG_SIZE)
pred = model.predict(tensor)[0]
idx = np.argmax(pred)
label = CLASS_NAMES[idx]
conf = pred[idx]
```

Discussion



Learning behavior

Training accuracy rises to ~ 0.99 , while validation peaks at ~ 0.90 then fluctuates. This is typical when model capacity is high relative to dataset size.



Model capacity hotspot

Flatten \rightarrow Dense(128) dominates parameters ($\sim 12.8M$). With limited images, the Dense layer can overfit. Using GlobalAveragePooling can cut parameters dramatically and often improves generalization.

Actionable improvements

- Early stopping on best val_acc (\approx Epoch 6)
- Replace Flatten with GlobalAveragePooling
- Stronger augmentation (color jitter) if lighting varies
- Add confusion matrix + per-class precision/recall
- Calibrate probabilities (temperature scaling) if needed

Limitations

Key limitations

- Dataset size and diversity: 900 images may not cover night scenes, rain, motion blur, or extreme angles.
- Background bias: model can learn correlated background cues (garage vs roadside) rather than tire shape.
- Ambiguity: flat vs full can be subtle from some viewpoints.
- Confidence ≠ calibrated probability: softmax scores can be overconfident without calibration.
- Real deployment needs robustness testing (cross-device/camera variation).

Future Work



Training & evaluation

- Early stopping + checkpointing
- Learning-rate scheduling
- Confusion matrix + macro F1
- Cross-validation (if data small)



Model improvements

- GlobalAveragePooling to reduce parameters
- Transfer learning (e.g., MobileNet/EfficientNet)
- Grad-CAM for explainability
- Quantization for edge devices

Deployment upgrades

- Save model in native Keras format (.keras)
- Add batch inference + caching for throughput
- Add input validation + OOD detection
- Track experiments (MLflow/W&B)
- Package as Docker for reproducible demos

Conclusion

- A simple 3-block CNN can classify tire conditions with strong validation performance (peak ~90%).
- Data augmentation + dropout help generalization on limited datasets.
- Streamlit provides an easy path to an interactive demo (upload → preprocess → predict).

Questions?

Thank you